
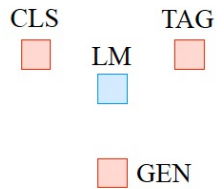
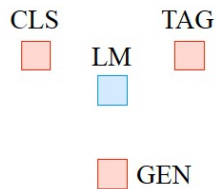
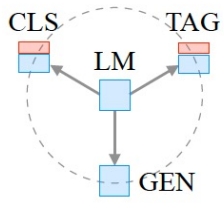
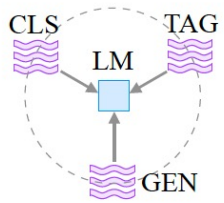


Prompting Methods in NLP

Qintong Li
2021/08/05



Four Paradigm in NLP

Paradigm	Engineering	Task Relation
a. Fully Supervised Learning (Non-Neural Network)	Features (e.g. word identity, part-of-speech, sentence length)	
b. Fully Supervised Learning (Neural Network)	Architecture (e.g. convolutional, recurrent, self-attentional)	
c. Pre-train, Fine-tune	Objective (e.g. masked language modeling, next sentence prediction)	
d. Pre-train, Prompt, Predict	Prompt (e.g. cloze, prefix)	

- Traditional supervised learning (using supervised dataset):

$$P(y|x; \theta)$$

- Prompt-based learning:

$$P(x; \theta)$$

- Prompt addition: $x' = f_{prompt}(x)$ [X] Overall, it was a [Z] movie.

“I love this movie. Overall it was a [Z] movie.”

- Answer search: $f_{fill}(x', z);$ I love this movie. Overall, it was a bad movie.

$$\hat{z} = \underset{z \in \mathcal{Z}}{\text{search}} P(f_{fill}(x', z); \theta). \quad \text{I love this movie. Overall, it was a good movie.}$$

- Answer mapping: $\hat{z} \rightarrow y$ the answer itself is the output
(e.g. “excellent”, “fabulous”, “wonderful”) to represent a single class (e.g. “++”)

Two advantages:

1. It allows the language model to be pre-trained on massive amounts of raw text;
2. By defining a new prompting function the model is able to perform few-shot or even zero-shot learning.

Design Considerations for Prompting

- Pre-trained Model Choice $P(x; \theta)$
- Prompt Engineering f_{prompt}
- Answer Engineering design \mathcal{Z}
- Expanding the Paradigm e.g., use of multiple prompts
- Prompt-based Training Strategies LM, Prompt

Design Considerations for Prompting

- Pre-trained Model Choice $P(x; \theta)$
- Prompt Engineering f_{prompt}
- Answer Engineering
- Expanding the Paradigm e.g., use of multiple prompts
- Prompt-based Training Strategies LM, Prompt

Prompt Engineering $f_{\text{prompt}}(x)$

- Prompt shape
 - Cloze prompts: fill in the blanks of a textual string
 - Prefix prompts: continue a string prefix
- Prompt acquisition
 - Manual template engineering
 - Automated template learning
 - Discrete prompts: the prompt is an actual text string
 - Continuous prompts: the prompt is instead described in the embedding space of the underlying LM

Discrete Prompts

AUTOPROMPT: Eliciting Knowledge from Language Models with Automatically Generated Prompts

Taylor Shin*[◇] **Yasaman Razeghi***[◇] **Robert L. Logan IV***[◇]

Eric Wallace[♠] **Sameer Singh**[◇]

[◇]University of California, Irvine [♠]University of California, Berkeley

{tshin1, yrazeghi, rlogan, sameer}@uci.edu

ericwallace@berkeley.edu

Motivations

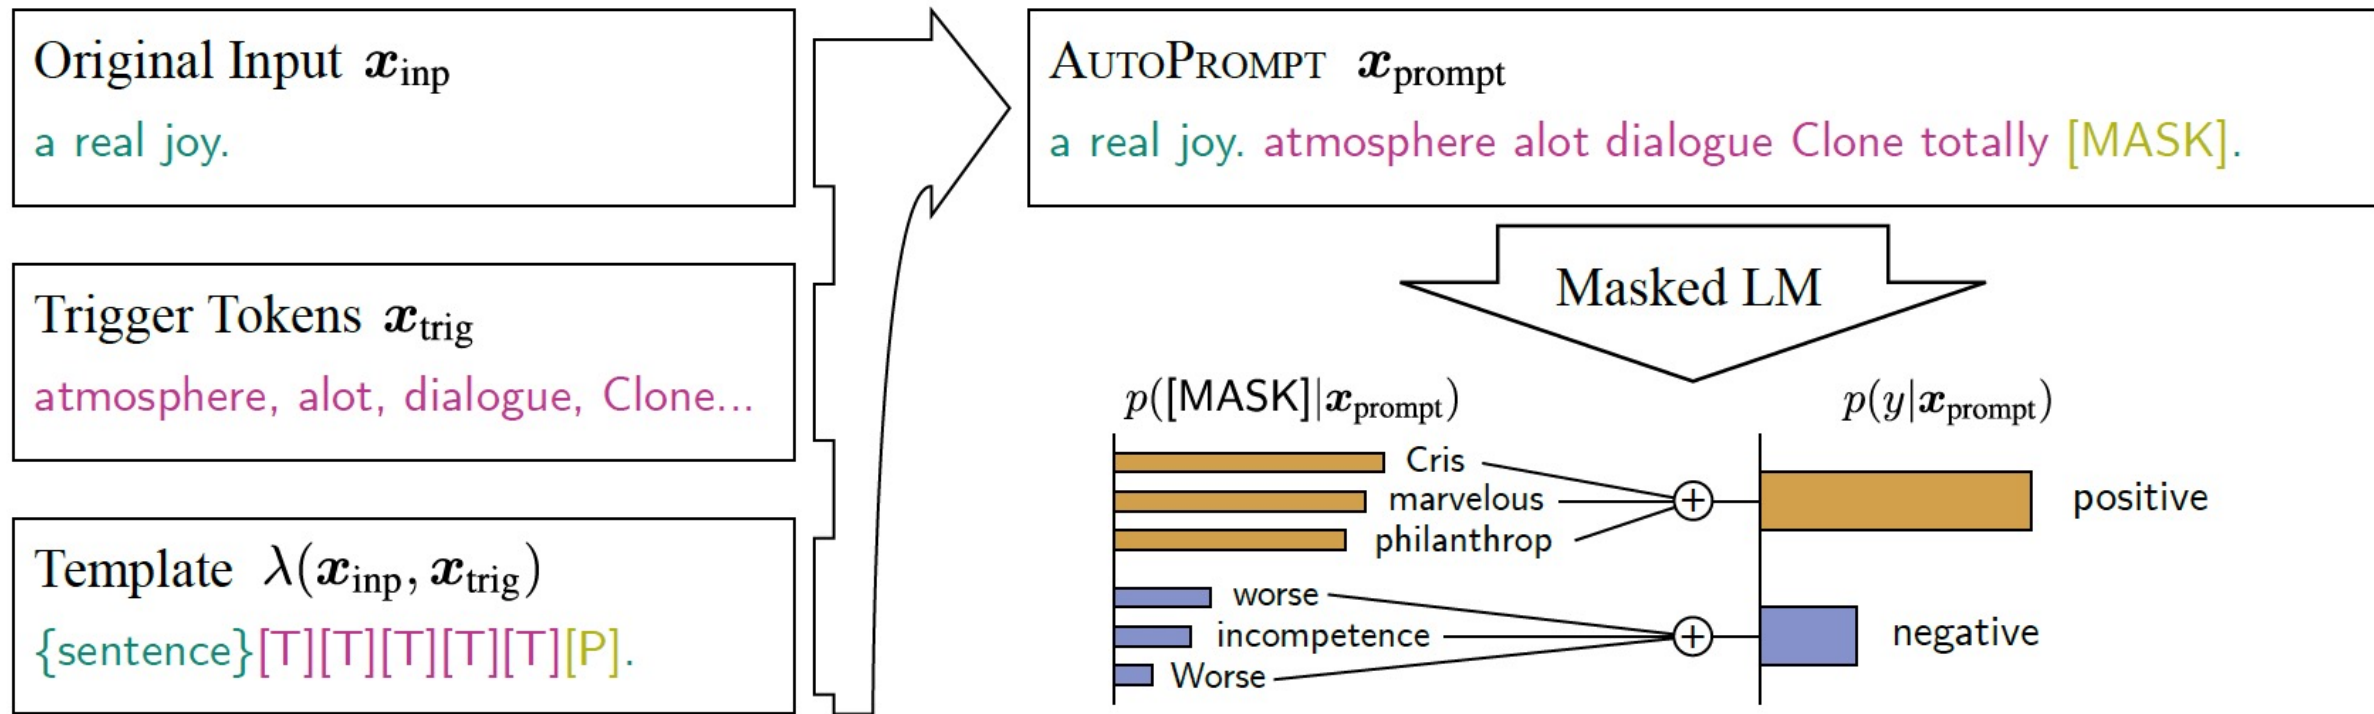
How can we directly evaluate the knowledge present in pretrained LMs, be it linguistic, factual, commonsense, or task-specific?

- Probing classifiers
 - false positives
 - high probing accuracy is not a sufficient condition
- Attention visualization
 - attention scores may be correlated with, but not caused by the underlying target knowledge
- Prompting
 - Provide a lower bound on what the language model “knows”

Contributions

AUTOPROMPT—an automated method for generating prompts: Given a task, e.g., sentiment analysis, AUTOPROMPT creates a prompt by combining the original task inputs (e.g. reviews) with a collection of trigger tokens according to a template.

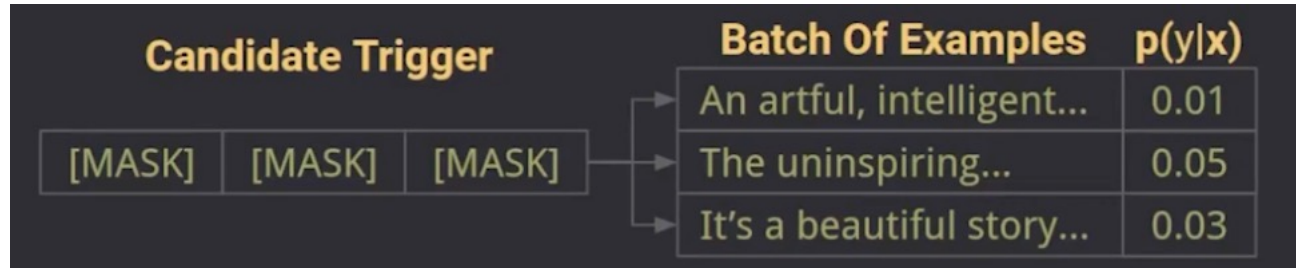
- Parameter-free
- Use output instead of internal representation
- Single model can solve many tasks
- No need manual effort
- Better than hand written and sometimes better than fine-tuning



Trigger Search

{sentence}[T][T][T][T][T][P].

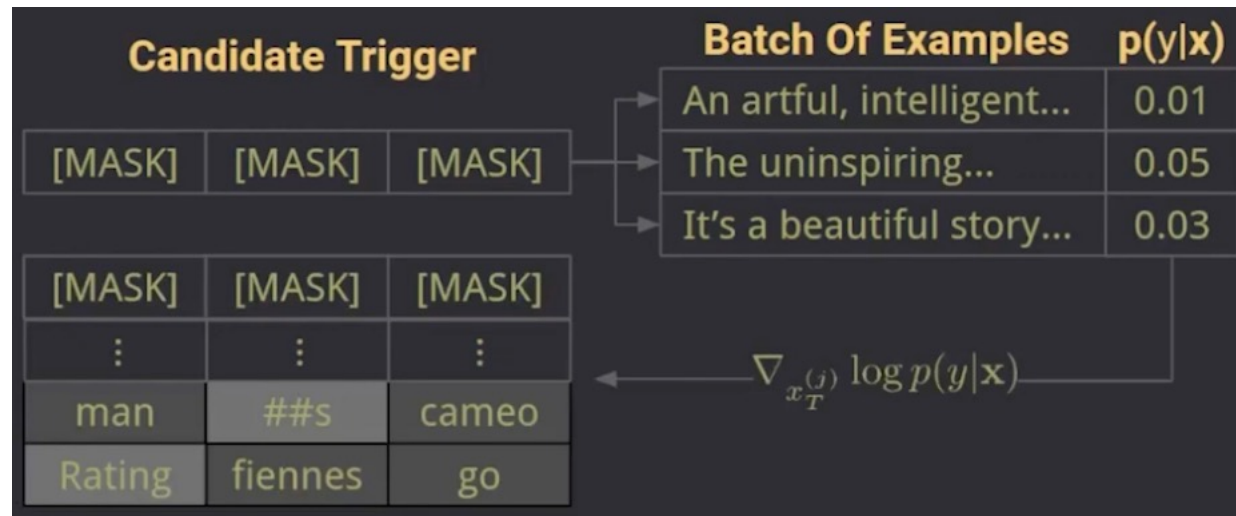
The idea is to add a number of “trigger” tokens that are shared across all prompts (denoted by [T]) to help model predict the correct label .



Trigger Search

{sentence}[T][T][T][T][T][P].

The idea is to add a number of “trigger” tokens that are shared across all prompts (denoted by [T]) to help model predict the correct label .

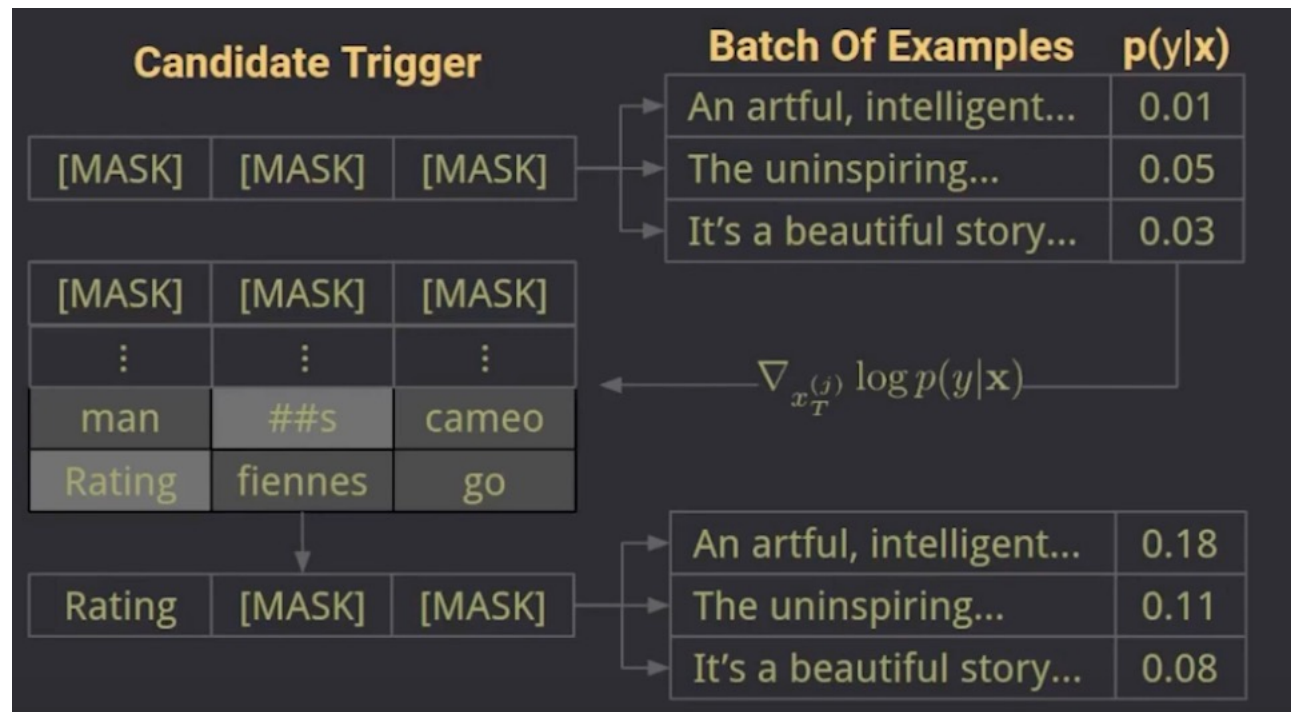


$$p(y|\mathbf{x}_{\text{prompt}}) = \sum_{w \in \mathcal{V}_y} p([\text{MASK}] = w | \mathbf{x}_{\text{prompt}})$$

Trigger Search

{sentence}[T][T][T][T][T][P].

The idea is to add a number of “trigger” tokens that are shared across all prompts (denoted by [T]) to help model predict the correct label .

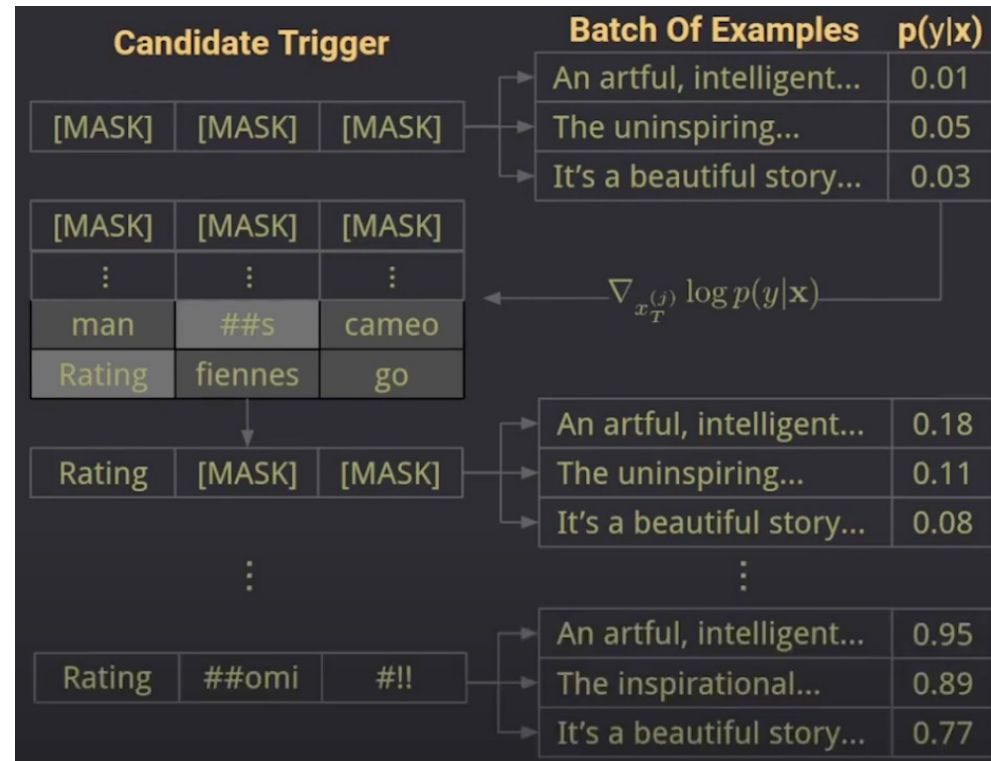


$$\mathcal{V}_{\text{cand}} = \text{top-}k_{w \in \mathcal{V}} [\mathbf{w}_{\text{in}}^T \nabla \log p(y|\mathbf{x}_{\text{prompt}})]$$

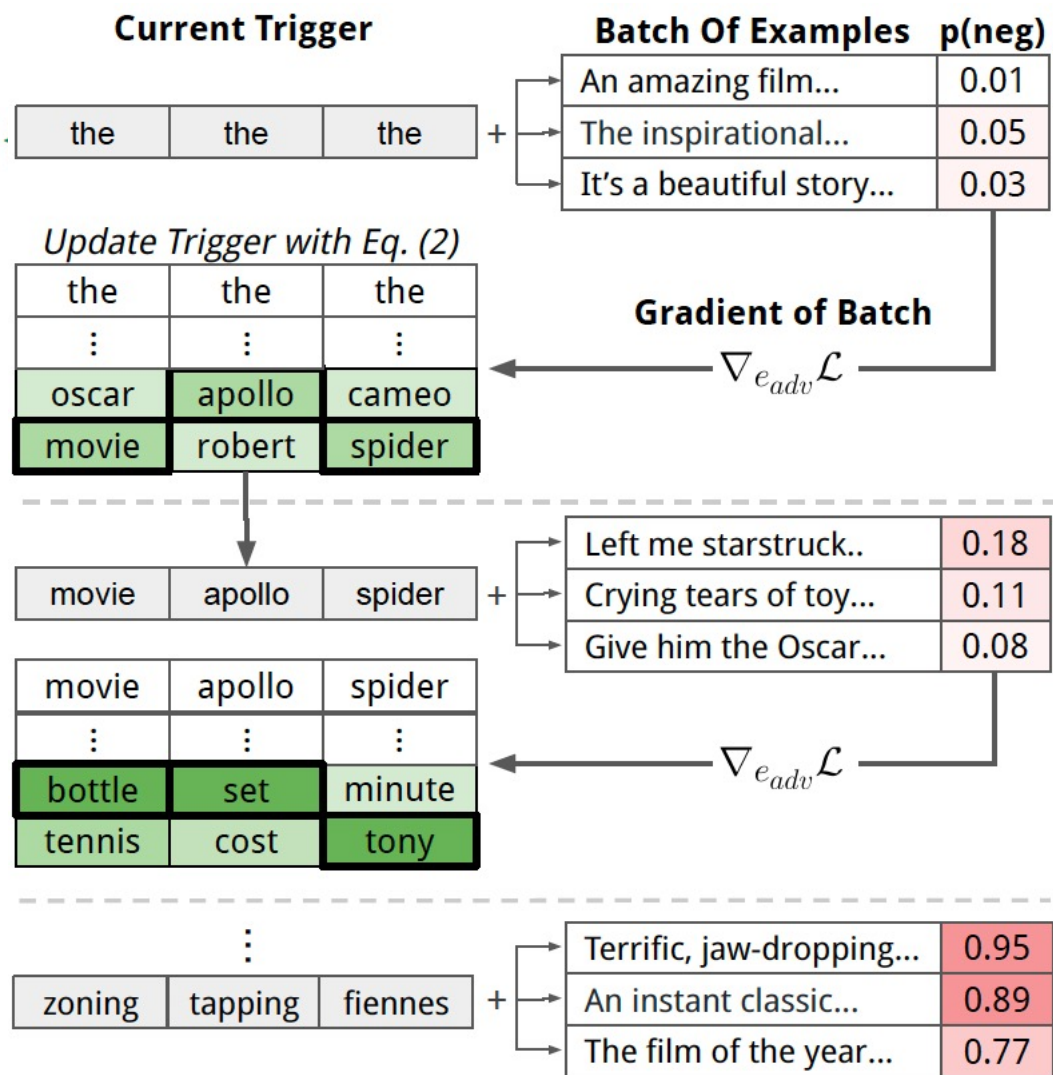
Trigger Search

{sentence}[T][T][T][T][T][P].

The idea is to add a number of “trigger” tokens that are shared across all prompts (denoted by [T]) to help model predict the correct label .




Trigger Search



Label Tokens

- Train a logistic classifier to predict the class label using the contextualized embedding of the [MASK] token as input.
- Find the words whose vector representations in the output layer of the language model are closest to the learned representations of these labels.



Experiments

Task	Prompt Template	Prompt found by AUTOPROMPT	Label Tokens
Sentiment Analysis	{sentence} [T]...[T] [P].	unflinchingly bleak and desperate Writing academicswhere overseas will appear [MASK].	pos: partnership, extraordinary, ##bla neg: worse, persisted, unconstitutional
NLI	{prem}[P][T]...[T]{hyp}	Two dogs are wrestling and hugging [MASK] concretepathic workplace There is no dog wrestling and hugging	con: Nobody, nobody, nor ent: ##found, ##ways, Agency neu: ##ponents, ##lary, ##uated
Fact Retrieval	<i>X plays Y music</i> {sub}[T]...[T][P].	Hall Overton fireplacemade antique son alto [MASK].	
Relation Extraction	<i>X is a Y by profession</i> {sent}{sub}[T]...[T][P].	Leonard Wood (born February 4, 1942) is a former Canadian politician. Leonard Wood gymnasium brotherdicative himself another [MASK].	

Table 3: **Example Prompts** by AUTOPROMPT for each task. On the left, we show the prompt template, which combines the input, a number of trigger tokens [T], and a prediction token [P]. For classification tasks (sentiment analysis and NLI), we make predictions by summing the model’s probability for a number of automatically selected label tokens. For fact retrieval and relation extraction, we take the most likely token predicted by the model.

Experiments

Manual prompt VS AutoPrompt

Model	Dev	Test
BiLSTM	-	82.8 [†]
BiLSTM + ELMo	-	89.3 [†]
BERT (linear probing)	85.2	83.4
BERT (finetuned)	-	93.5 [†]
RoBERTa (linear probing)	87.9	88.8
RoBERTa (finetuned)	-	96.7 [†]
BERT (manual)	63.2	63.2
BERT (AUTOPROMPT)	80.9	82.3
RoBERTa (manual)	85.3	85.2
RoBERTa (AUTOPROMPT)	91.2	91.4


Prompt Type	Original			T-REx		
	MRR	P@10	P@1	MRR	P@10	P@1
LAMA	40.27	59.49	31.10	35.79	54.29	26.38
LPAQA (Top1)	43.57	62.03	34.10	39.86	57.27	31.16
AUTOPROMPT 5 Tokens	53.06	72.17	42.94	54.42	70.80	45.40
AUTOPROMPT 7 Tokens	53.89	73.93	43.34	54.89	72.02	45.57

Factual Retrieval


Table 1: **Sentiment Analysis** performance on the SST-2 test set of supervised classifiers (top) and fill-in-the-blank MLMs (bottom). Scores marked with [†] are from the GLUE leaderboard: <http://gluebenchmark.com/leaderboard>.

Discussion


- Prompting as an Alternative to Finetuning
 - Prompts generated using AUTOPROMPT can achieve higher accuracy than finetuning in the low-data regime.
 - Prompting has advantages over finetuning when trying to solve many different tasks.




(a) BERT on SST-2



(b) RoBERTa on SST-2



(c) BERT on SICK-E



(d) RoBERTa on SICK-E

Discussion

- Limitations of Prompting
 - Preliminary evaluation on datasets such as QQP (Iyer et al., 2017) and RTE (Dagan et al., 2005), prompts generated manually and with AUTOPROMPT did not perform considerably better than chance.
 - AUTOPROMPT makes prompt-based probes more generally applicable, but, it still remains just one tool in the toolbox of the interpretability researcher.

Discussion

- Limitations of AUTOPROMPT
 - It requires labeled training data.
 - AUTOPROMPT generated prompts lack interpretability.
 - It can sometimes struggle when the training data is highly imbalanced.
 - The large discrete space of phrases (need more more effective crafting techniques).

Model	SICK-E Datasets		
	standard	3-way	2-way
Majority	56.7	33.3	50.0
BERT (finetuned)	86.7	84.0	95.6
BERT (linear probing)	68.0	49.5	91.9
RoBERTa (linear probing)	72.6	49.4	91.1
BERT (AUTOPROMPT)	62.3	55.4	85.7
RoBERTa (AUTOPROMPT)	65.0	69.3	87.3

Prompt-based Training Strategies

Strategy	LM Params	Prompt Params		Example
		Additional	Tuned	
Promptless Fine-tuning	Tuned	-		ELMo [130], BERT [32], BART [94]
Tuning-free Prompting	Frozen	✗	✗	GPT-3 [16], AutoPrompt [159], LAMA [133]
Fixed-LM Prompt Tuning	Frozen	✓	Tuned	Prefix-Tuning [96], Prompt-Tuning [91]
Fixed-prompt LM Tuning	Tuned	✗	✗	PET-TC [153], PET-Gen [152], LM-BFF [46]
Prompt+LM Fine-tuning	Tuned	✓	Tuned	PADA [8], P-Tuning [103], PTR [56]

Table 6: Characteristics of different tuning strategies. “Additional” represents if there are additional parameters beyond LM parameters while “Tuned” denotes if parameters are updated.

Prefix-Tuning: Optimizing Continuous Prompts for Generation

Xiang Lisa Li

Stanford University

`xlisali@stanford.edu`

Percy Liang

Stanford University

`pliang@cs.stanford.edu`

Overview

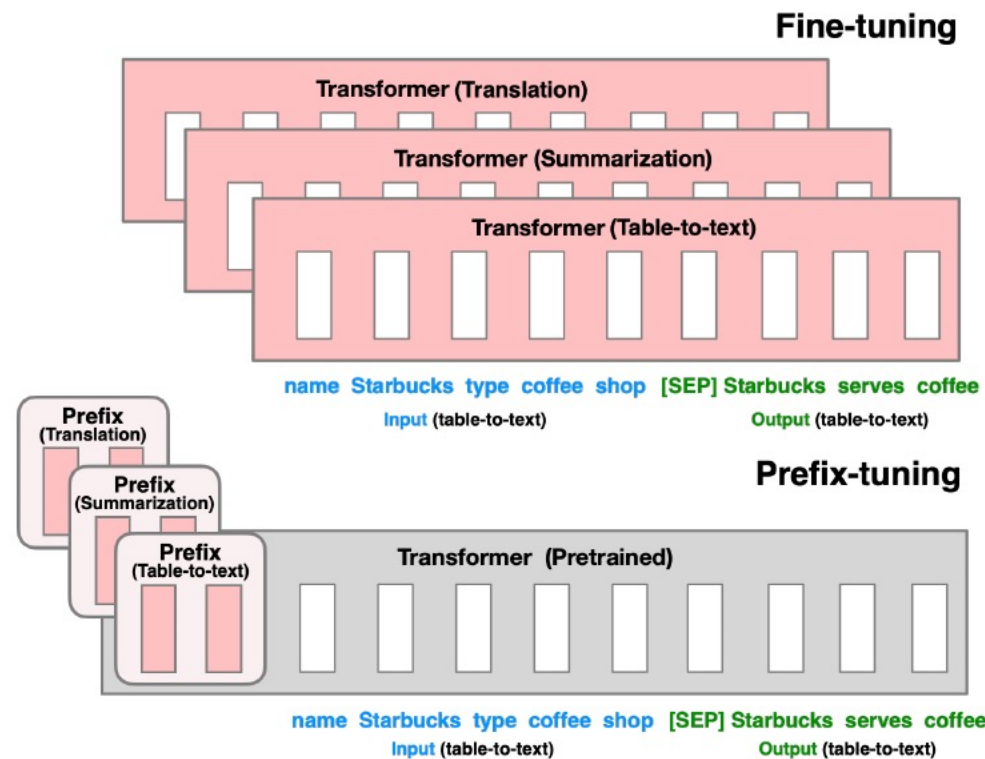


Figure 1: Fine-tuning (top) updates all Transformer parameters (the red Transformer box) and requires storing a full model copy for each task. We propose prefix-tuning (bottom), which freezes the Transformer parameters and only optimizes the prefix (the red prefix blocks). Consequently, we only need to store the prefix for each task, making prefix-tuning modular and space-efficient. Note that each vertical block denote transformer activations at one time step.

Method

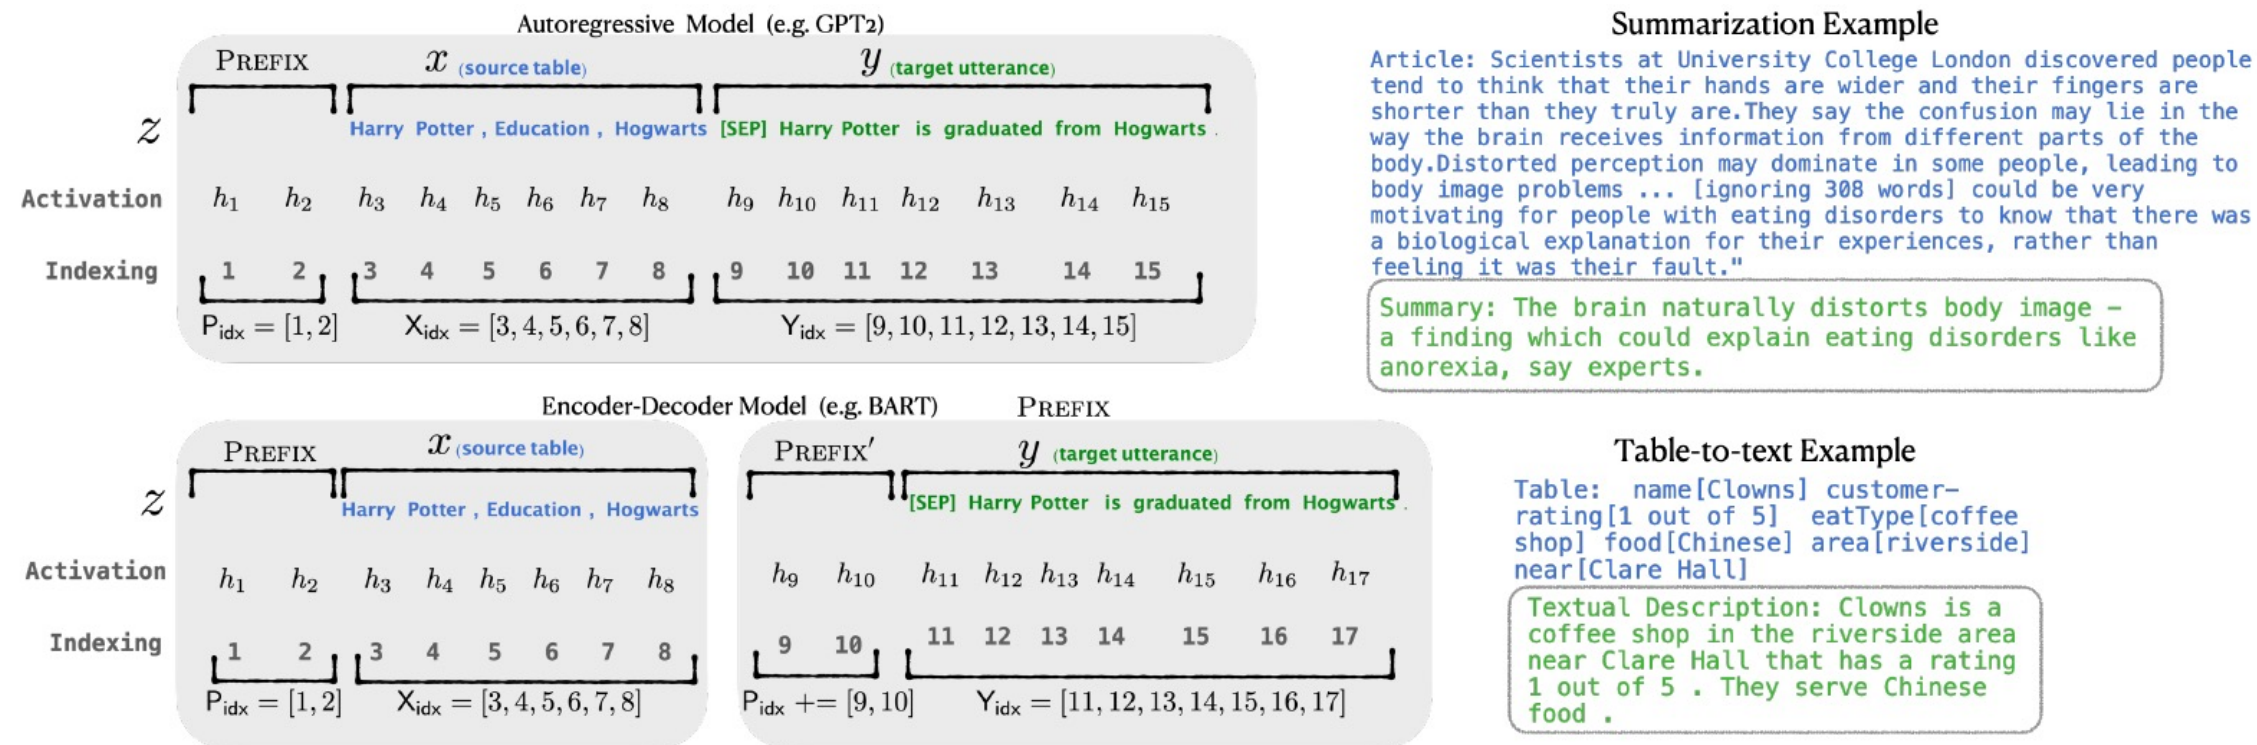


Figure 2: An annotated example of prefix-tuning using an autoregressive LM (top) and an encoder-decoder model (bottom). The prefix activations $\forall i \in P_{idx}, h_i$ are drawn from a trainable matrix P_θ . The remaining activations are computed by the Transformer.

Experiments

	E2E					WebNLG									DART					
	BLEU	NIST	MET	R-L	CIDE _r	BLEU			MET			TER ↓			BLEU	MET	TER ↓	Mover	BERT	BLEURT
						S	U	A	S	U	A	S	U	A						
GPT-2 _{MEDIUM}																				
FINE-TUNE	68.2	8.62	46.2	71.0	2.47	64.2	27.7	46.5	0.45	0.30	0.38	0.33	0.76	0.53	46.2	0.39	0.46	0.50	0.94	0.39
FT-TOP2	68.1	8.59	46.0	70.8	2.41	53.6	18.9	36.0	0.38	0.23	0.31	0.49	0.99	0.72	41.0	0.34	0.56	0.43	0.93	0.21
ADAPTER(3%)	68.9	8.71	46.1	71.3	2.47	60.4	48.3	54.9	0.43	0.38	0.41	0.35	0.45	0.39	45.2	0.38	0.46	0.50	0.94	0.39
ADAPTER(0.1%)	66.3	8.41	45.0	69.8	2.40	54.5	45.1	50.2	0.39	0.36	0.38	0.40	0.46	0.43	42.4	0.36	0.48	0.47	0.94	0.33
PREFIX(0.1%)	69.7	8.81	46.1	71.4	2.49	62.9	45.6	55.1	0.44	0.38	0.41	0.35	0.49	0.41	46.4	0.38	0.46	0.50	0.94	0.39
GPT-2 _{LARGE}																				
FINE-TUNE	68.5	8.78	46.0	69.9	2.45	65.3	43.1	55.5	0.46	0.38	0.42	0.33	0.53	0.42	47.0	0.39	0.46	0.51	0.94	0.40
Prefix	70.3	8.85	46.2	71.7	2.47	63.4	47.7	56.3	0.45	0.39	0.42	0.34	0.48	0.40	46.7	0.39	0.45	0.51	0.94	0.40
SOTA	68.6	8.70	45.3	70.8	2.37	63.9	52.8	57.1	0.46	0.41	0.44	-	-	-	-	-	-	-	-	-

	R-1 ↑	R-2 ↑	R-L ↑
FINE-TUNE(Lewis et al., 2020)	45.14	22.27	37.25
PREFIX(2%)	43.80	20.93	36.05
PREFIX(0.1%)	42.92	20.03	35.05

Discussion

- Prefix-tuning, a lightweight alternative to fine-tuning that prepends a trainable continuous prefix for NLG tasks.
- Despite learning 1000x fewer parameters than finetuning, prefix-tuning can maintain a comparable performance in a full data setting and outperforms fine-tuning in both low-data and extrapolation settings.